

# CUDA SHARED MEMORY

NVIDIA Corporation

- Difference between *host* and *device*

*Host*                      CPU

*Device*                   GPU

- Using `__global__` to declare a function as device code
  - Executes on the device
  - Called from the host (or possibly from other device code)
- Passing parameters from host code to a device function


- Basic device memory management
  - `cudaMalloc()`
  - `cudaMemcpy()`
  - `cudaFree()`
- Launching parallel kernels
  - Launch **N** copies of `add()` with `add<<<N, 1>>>>(...)` ;
  - Use `blockIdx.x` to access block index

- Shared memory is equivalent to a user-managed cache: The application explicitly allocates and accesses it.


Recall: how does cache work in OS and computer organization courses

- Launching parallel kernels
  - Launch **N** copies of **add()** with **add<<<N, 1>>> (...)** ;
  - Use **blockIdx.x** to access block index

- Consider applying a 1D stencil to a 1D array of elements
  - Each output element is the sum of input elements within a radius
- If radius is 3, then each output element is the sum of 7 input elements:



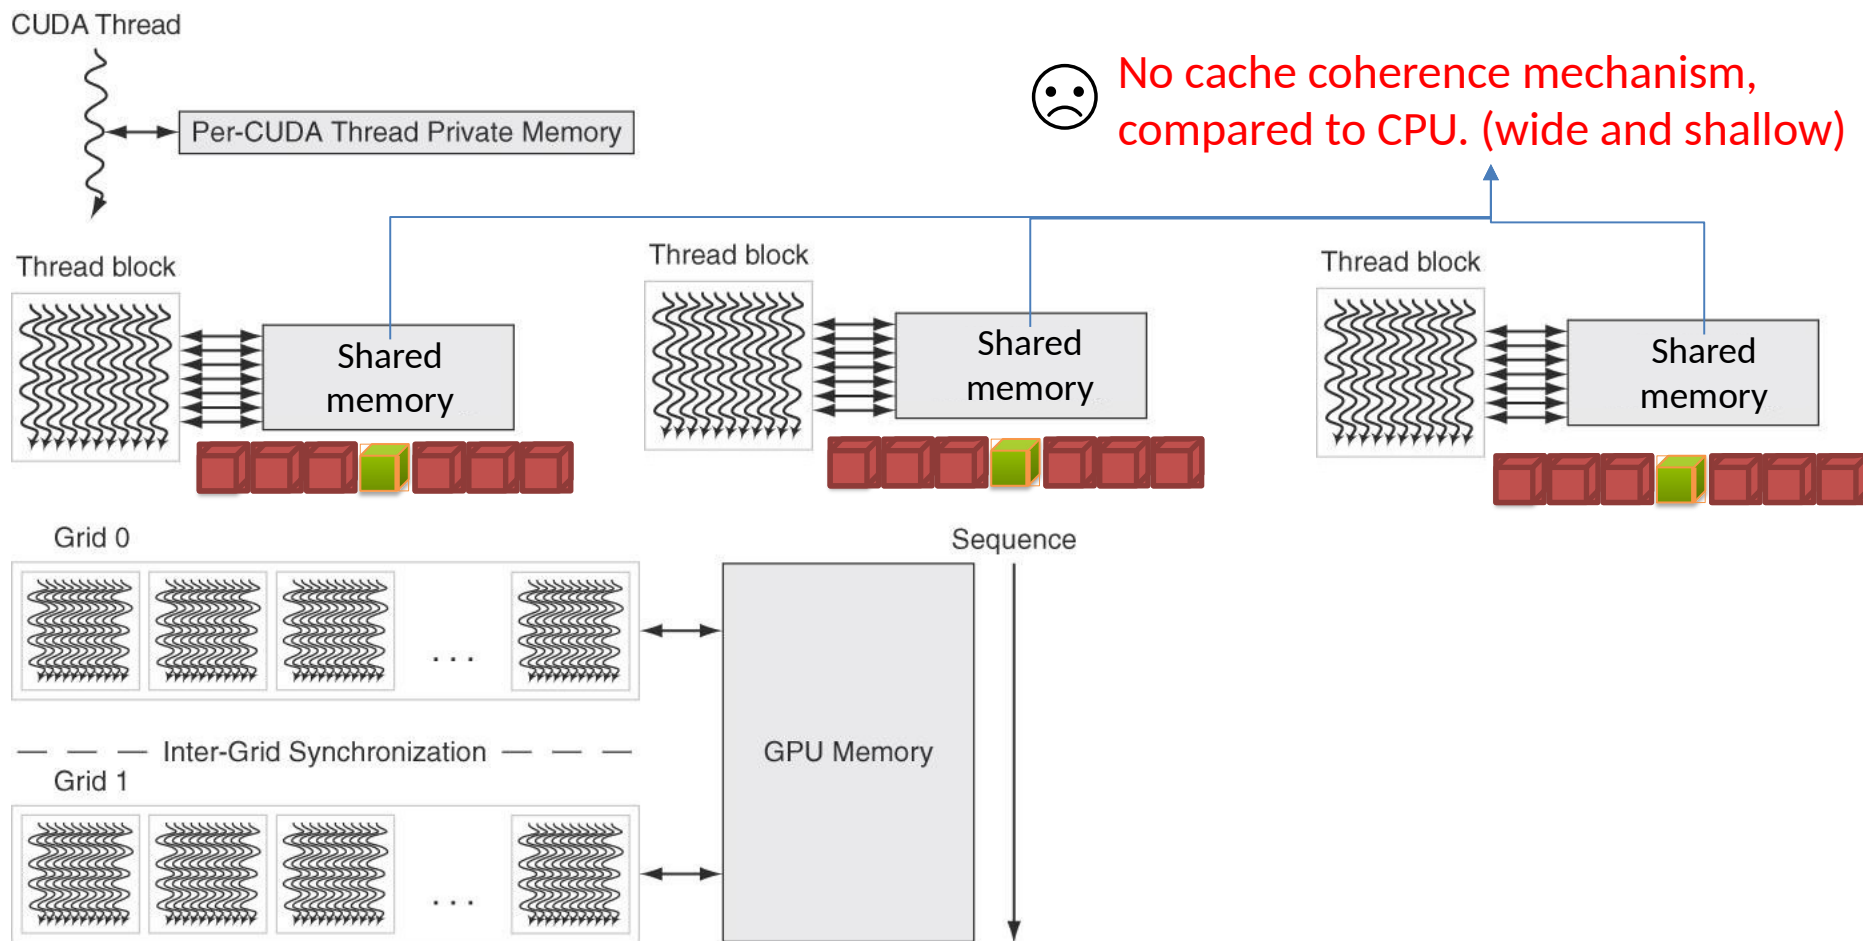
- Each thread processes one output element
- **blockDim.x** elements per block
- Input elements are read several times
  - With radius 3, each input element is read seven times




- Terminology: within a block, threads share data via **shared memory**
- Shared memory is equivalent to a user-managed cache:
  - The application explicitly allocates and accesses it.
- Extremely fast on-chip memory, user-managed
- Declare using **`__shared__`**, allocated a variable per block that:
  - Resides in the shared memory space of a thread block,
  - Has the lifetime of the block,
  - Has a distinct object per block,
  - Is only accessible from all the CUDA threads within the block,

- Typical programming pattern:
  - Load data from device memory to shared memory
  - Synchronize with all the other threads of the block so that each thread can safely read shared memory locations that were populated by different threads,
  - Process the data in shared memory
  - Synchronize again if necessary to make sure that shared memory has been updated with the results,
  - Write the results back to device memory.





- Cache data in shared memory
  - Read  $(\text{blockDim.x} + 2 * \text{radius})$  input elements from global memory to shared memory
  - Compute  $\text{blockDim.x}$  output elements
  - Write  $\text{blockDim.x}$  output elements to global memory
- Each block needs a halo of  $\text{radius}$  elements at each boundary




```
__global__ void stencil_1d(int *in, int *out) {  
    __shared__ int temp[BLOCK_SIZE + 2 * RADIUS];  
    int gindex = threadIdx.x + blockIdx.x * blockDim.x;  
    int lindex = threadIdx.x + RADIUS;
```

```
    // Read input elements into shared memory
```

```
    temp[lindex] = in[gindex];  
    if (threadIdx.x < RADIUS) {  
        temp[lindex - RADIUS] = in[gindex - RADIUS];  
        temp[lindex + BLOCK_SIZE] = in[gindex + BLOCK_SIZE];  
    }
```



A thread block[0-15] calculates stencil



The 3 CUDA threads move left 3  
orange elements from in[] to temp[]



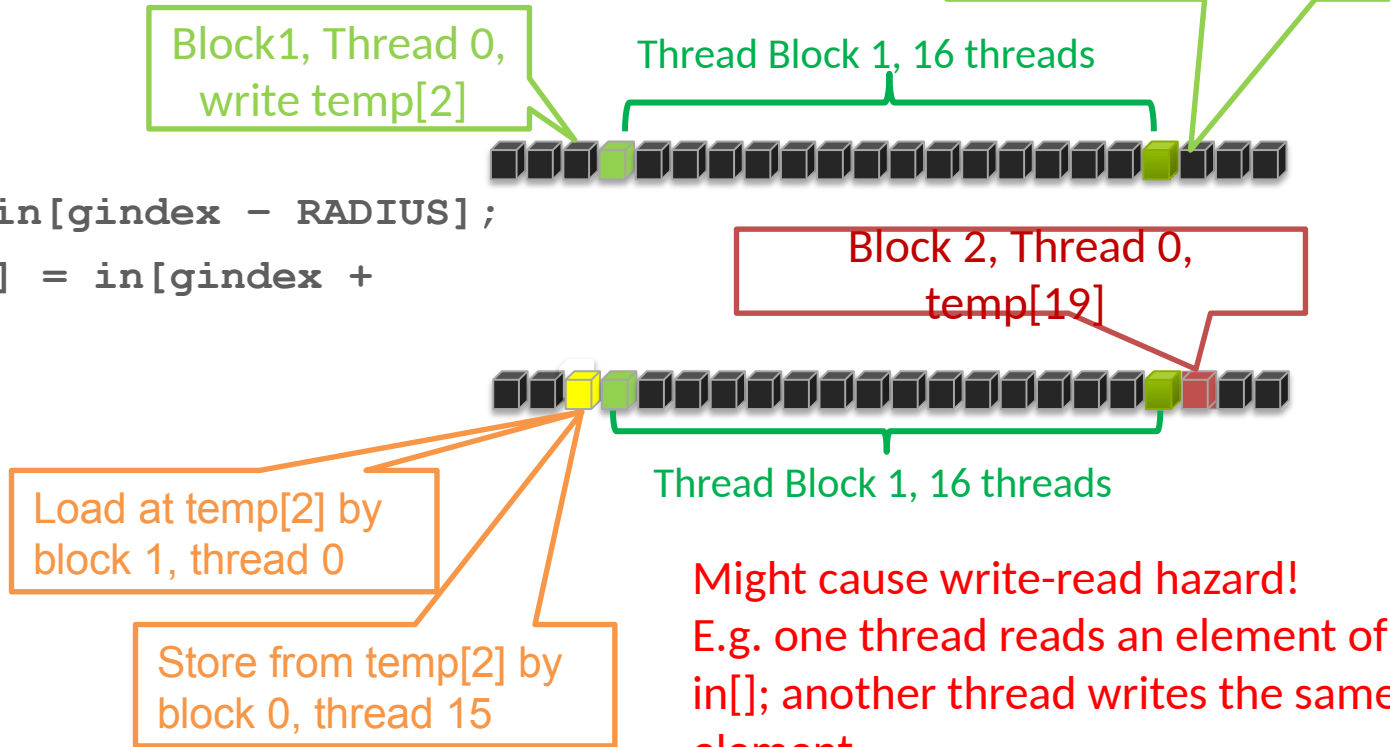
The 3 CUDA threads move right 3  
orange elements from in[] to temp[]



```
// Apply the stencil in shared memory  
int result = 0;  
for (int offset = -RADIUS ; offset <= RADIUS ; offset++)  
    result += temp[lindex + offset];  
  
// Store the result  
out[gindex] = result;  
}
```

- The stencil example will not work...
- Suppose **thread 15 in Block 1** reads the `temp[19]` for calculating stencil after **thread 0 in block 2** has fetched `temp[19]` from global memory `in[]`.

```
temp[lindex] = in[gindex];  
if (threadIdx.x < RADIUS) {  
    temp[lindex - RADIUS] = in[gindex - RADIUS];  
    temp[lindex + BLOCK_SIZE] = in[gindex +  
        BLOCK_SIZE];  
}  
int result = 0;  
result += temp[lindex + 1];
```



Might cause write-read hazard!  
E.g. one thread reads an element of `in[]`; another thread writes the same element.

- `void _syncthreads () ;`
- Synchronizes all threads within a CUDA block, how to explain in GPU architecture
  - Used to prevent RAW / WAR / WAW hazards
- All threads must reach the barrier
  - In conditional code, the condition must be uniform across the block

```
__global__ void stencil_1d(int *in, int *out) {
    __shared__ int temp[BLOCK_SIZE + 2 * RADIUS];
    int gindex = threadIdx.x + blockIdx.x * blockDim.x;
    int lindex = threadIdx.x + radius;

    // Read input elements into shared memory
    temp[lindex] = in[gindex];
    if (threadIdx.x < RADIUS) {
        temp[lindex - RADIUS] =
            in[gindex - RADIUS];
        temp[lindex +
            BLOCK_SIZE] = in[gindex
            + BLOCK_SIZE];
    }

    // Synchronize (ensure all
```

```
// Apply the stencil
int result = 0;
for (int offset = -RADIUS ; offset <= RADIUS ; offset++)
    result += temp[lindex + offset];

// Store the result
out[gindex] = result;
}
```



- Use `__shared__` to declare a variable/array in shared memory
  - Data is shared between threads in a block
  - Not visible to threads in other blocks
- Use `__syncthreads()` as a barrier
  - Use to prevent data hazards

- CUDA GPU architecture and basic optimizations
- Atomics, Reductions, Warp Shuffle
- Using Managed Memory
- Concurrency (streams, copy/compute overlap, multi-GPU)
- Analysis Driven Optimization
- Cooperative Groups

- Shared memory:
  - <https://devblogs.nvidia.com/using-shared-memory-cuda-cc/>
- CUDA Programming Guide:
  - <https://docs.nvidia.com/cuda/cuda-c-programming-guide/index.html#shared-memory>
- CUDA Documentation:
  - <https://docs.nvidia.com/cuda/index.html>
  - <https://docs.nvidia.com/cuda/cuda-runtime-api/index.html> (runtime API)