

ATOMICS, REDUCTIONS, WARP SHUFFLE

Bob Crovella, 5/13/2020

AGENDA

- Transformations vs. Reductions, Thread Strategy
- Atomics, Atomic Reductions
- Atomic Tips and Tricks
- Classical Parallel Reduction
- Parallel Reduction + Atomics
- Warp Shuffle, Reduction with Warp Shuffle
- Other Warp Shuffle Uses
- Further Study
- Homework


ATOMICS

Sum - reduction

```
const int size = 100000;  
float a[size] = {...};  
float sum = 0;  
for (int i = 0; i < size; i++)  
    sum += a[i];
```

-> sum variable contains the sum
of all the elements of array **a**


May guide the thread strategy: what will each thread do?



Transformation:

e.g. $c[i] = a[i] + 10;$

CUDA thread strategy: one CUDA thread per output point



Reduction:

e.g. $*c = \sum a[i]$

CUDA thread strategy: ??

One thread per input point

```
*c += a[i];
```

(Doesn't work.) Actual code the GPU executes:

LD R2, a[i] (Thread independent)

LD R1, c (READ)

ADD R3, R1, R2 (MODIFY)

ST c, R3 (WRITE)

But **every CUDA Thread** is trying to do this, potentially at the same time

The CUDA programming model does not enforce any order of thread execution

ATOMICS TO THE RESCUE

indivisible READ-MODIFY-WRITE

`old = atomicAdd (&c, a[i]);` <https://docs.nvidia.com/cuda/cuda-c-programming-guide/index.html#atomic-functions>

Reads the 16-bit, 32-bit or 64-bit word `old` located at the `address` in global or shared memory, computes (`old` + val), and stores the result back to memory at the same `address`. These three operations are performed in one atomic transaction. The function returns `old`.

LD R2, a[i] (CUDA thread independent)

LD R1, c	(READ)
ADD R3, R1, R2	(MODIFY)
ST R3, c	(WRITE)

Becomes one indivisible operation/instruction:

→ **RED.E.ADD.F32.FTZ.RN [c], R2;**

Facilitated by special hardware in the L2 cache

May have performance implications

atomicAdd implementation

// Note that any atomic operation can be implemented based on atomicCAS() (Compare And Swap). For example.....

```
#if __CUDA_ARCH__ < 600
__device__ double atomicAdd(double* address, double val)
{
    unsigned long long int* address_as_ull = (unsigned long long int*)address;
    unsigned long long int old = *address_as_ull, assumed;

    do {
        assumed = old;
        old = atomicCAS(address_as_ull, assumed,
            __double_as_longlong(val + __longlong_as_double(assumed)));

        // Note: uses integer comparison to avoid hang in case of NaN (since NaN != NaN)
    } while (assumed != old);

    return __longlong_as_double(old);
}
#endif
```

atomicCAS(int* address, int compare, int val):

Reads the 16-bit, 32-bit or 64-bit word **old** located at the **address** in global or shared memory, computes (old == compare ? val : old) , and stores the result back to memory at the same **address**. These three operations are performed in one atomic transaction. The function returns **old** (Compare And Swap).

- atomicMax/Min – choose the max (or min)
- atomicAdd/Sub – add to (or subtract from)
- atomicInc/Dec – increment (or decrement) and account for rollover/underflow
- atomicExch/CAS – swap values, or conditionally swap values
- atomicAnd/Or/Xor – bitwise ops
- atomics have different datatypes they can work on (e.g. int, unsigned, float, etc.)
- <https://docs.nvidia.com/cuda/cuda-c-programming-guide/index.html#atomic-functions>
- **Reference textbook “The Art of Multiprocessor Programming” by MIT**


Determine my place in an order

- Could be used to determine next work item, queue slot, etc.
- `int my_position = atomicAdd(&order, 1);`
- Most atomics return a value that is the “old” value that was in the location receiving the atomic update.

CLASSICAL PARALLEL REDUCTION

Parallel Reduction

Tree-based approach used within each thread block



Need to be able to use multiple thread blocks

- To process very large arrays

- To keep all streaming multiprocessors (SIMD processors) on the GPU busy


- Each thread block reduces a portion of the array

But how do we communicate partial results between thread blocks?

Problem: Global Synchronization

- **If we could synchronize across all thread blocks, could easily reduce very large arrays, right?**
 - Global sync after each block produces its result
 - Once all blocks reach sync, continue recursively
- **But CUDA has no global synchronization. Why?**
 - Expensive to build in hardware for GPUs with high processor count
 - Would force programmer to run fewer blocks to avoid deadlock, which may reduce overall efficiency
- **Solution: decompose into multiple kernels**
 - Kernel launch serves as a global synchronization point
 - Kernel launch has HW overhead and SW overhead

Avoid global sync by decomposing computation into multiple kernel invocations



In the case of reductions, code for all levels is the same

Recursive kernel invocation

Strive to reach GPU peak performance

Choose the right metric:

GFLOP/s: for compute-bound kernels Bandwidth: for
memory-bound kernels

Reductions have very low arithmetic intensity

1 flop per element loaded (bandwidth-optimal), limited by bandwidth

Therefore we should strive for peak bandwidth

Will use G80 GPU for this example

384-bit memory interface, 900 MHz DDR


$$384 * 800 * 2 / 8 = 86.4 \text{ GB/s}$$

V100 GPU memory bandwidth 900 GB/s

```
global __void reduce0(int *g_idata, int *g_odata) {  
    // g_odata's size is the number of blocks,  
    // g_idata's size <= number of CUDA threads in grid.  
    __shared__ int sdata[BLOCK_SIZE];  
    // each thread loads one element from global to shared mem  
    unsigned int tid = threadIdx.x;  
    unsigned int i = blockIdx.x*blockDim.x + threadIdx.x;  
    sdata[tid] = g_idata[i];  
    syncthreads();  
    // do reduction in shared mem  
    for(unsigned int s=1; s < blockDim.x; s *= 2) {  
        if (tid % (2*s) == 0) {  
            sdata[tid] += sdata[tid + s];  
        }  
        syncthreads();  
    }  
    // write result for this block to global mem  
    if (tid == 0) g_odata[blockIdx.x] = sdata[0];  
}
```


Parallel Reduction: Interleaved Addressing

Warp 0 Thread ID: 0, 2, 4, 6,, 32, participate computing



```
global_ void reduce1(int *g_idata, int *g_odata)
{ __shared__ int sdata[BLOCK_SIZE];
```

```
// each thread loads one element from global to shared mem
```

```
    unsigned int tid = threadIdx.x;
    unsigned int i = blockIdx.x*blockDim.x + threadIdx.x;
    sdata[tid] = g_idata[i];
    syncthreads();
```

```
// do reduction in shared mem
```

```
for (unsigned int s=1; s < blockDim.x; s *= 2) {
    if (tid % (2*s) == 0) {
        sdata[tid] += sdata[tid + s];
    }
    syncthreads();
}
```

Problem: highly divergent warps are very inefficient, and % operator is very slow

```
// write result for this block to global mem
```

```
if (tid == 0) g_odata[blockIdx.x] = sdata[0];
}
```

	Time (2^{22} ints)	Bandwidth
Kernel 1: interleaved addressing with divergent branching	8.054 ms	2.083 GB/s

Note: Block Size = 128 threads for all tests

Reduction #2: Interleaved Addressing

Just replace **warp divergent branch** in loop:

```
for (unsigned int s=1; s < blockDim.x; s *= 2)
    { if (tid % (2*s) == 0) {
        sdata[tid] += sdata[tid + s];
    }
    -   syncthreads();
}
```



With strided index and **warp non-divergent branch**:

```
for (unsigned int s=1; s < blockDim.x; s *= 2)    {
    int index = 2 * s * tid;

    if (index < blockDim.x) {
        // warps are non-divergent, either a warp of 32
        // threads doing if code block, or skipping if.
        sdata[index] += sdata[index + s];
    }
    -   syncthreads();
}
```

Parallel Reduction: Interleaved Addressing

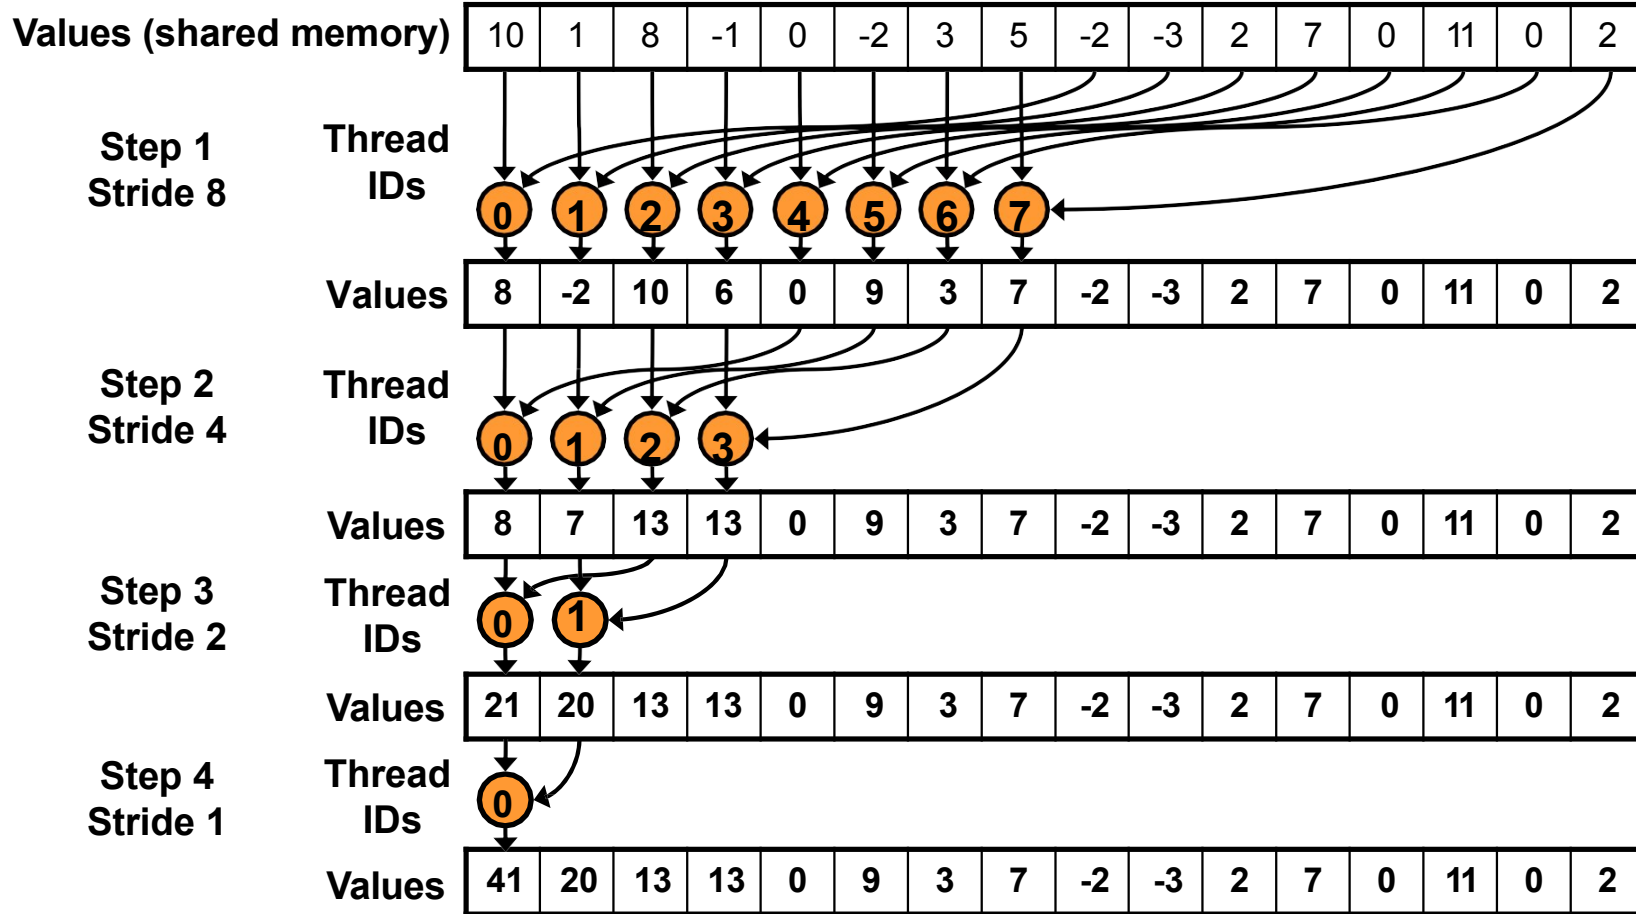
Warp 0 Thread ID: 0, 1, 2, 3, 4, 5, 6,, 31, participate computing



New Problem: Shared Memory Bank Conflicts

Performance for 4M element reduction

	Time (2^{22} ints)	Bandwidth	Step Speedup	Cumulative Speedup
Kernel 1: interleaved addressing with divergent branching	8.054 ms	2.083 GB/s		
Kernel 2: interleaved addressing with bank conflicts	3.456 ms	4.854 GB/s	2.33x	2.33x



Sequential addressing is conflict free

Reduction #3: Sequential Addressing

Just replace strided indexing in inner loop:

```
for (unsigned int s=1; s < blockDim.x; s *= 2)
{
    // s represents stride
    int index = 2 * s * tid;

    if (index < blockDim.x) {
        sdata[index] += sdata[index + s];
    }

    syncthreads();
}
```

With reversed loop and threadID-based indexing:

```
for (unsigned int s=blockDim.x/2; s>0; s>>=1) {
    // s represents stride
    if (tid < s) {
        sdata[tid] += sdata[tid + s];
    }

    syncthreads();
}
```


Performance for 4M element reduction

	Time (2^{22} ints)	Bandwidth	Step Speedup	Cumulative Speedup
Kernel 1: interleaved addressing with divergent branching	8.054 ms	2.083 GB/s		
Kernel 2: interleaved addressing with bank conflicts	3.456 ms	4.854 GB/s	2.33x	2.33x
Kernel 3: sequential addressing	1.722 ms	9.741 GB/s	2.01x	4.68x

Problem:

```
for (unsigned int s=blockDim.x/2; s>0; s>>=1) {  
    if (tid < s) {  
        sdata[tid] += sdata[tid + s];  
    }  
    syncthreads();  
}
```

Half of the threads are idle on first loop iteration!

This is wasteful...

Reduction #4: First Add During Load

Halve the number of blocks, and replace single load:

```
// each thread loads one element from global to shared mem
unsigned int tid = threadIdx.x;
unsigned int i = blockIdx.x*blockDim.x + threadIdx.x;
sdata[tid] = g_idata[i];
syncthreads();
```

With two loads and first add of the reduction:

```
// perform first level of reduction,
// reading from global memory, writing to shared memory
unsigned int tid = threadIdx.x;
unsigned int i = blockIdx.x*(blockDim.x*2) + threadIdx.x;
sdata[tid] = g_idata[i] + g_idata[i+blockDim.x];
syncthreads();
```

Performance for 4M element reduction

	Time (2^{22} ints)	Bandwidth	Step Speedup	Cumulative Speedup
Kernel 1: interleaved addressing with divergent branching	8.054 ms	2.083 GB/s		
Kernel 2: interleaved addressing with bank conflicts	3.456 ms	4.854 GB/s	2.33x	2.33x
Kernel 3: sequential addressing	1.722 ms	9.741 GB/s	2.01x	4.68x
Kernel 4: first add during global load	0.965 ms	17.377 GB/s	1.78x	8.34x

Instruction Bottleneck

At 17 GB/s, we're far from bandwidth bound

And we know reduction has **low arithmetic intensity**

→ Therefore a likely bottleneck is instruction overhead
Ancillary instructions that are not loads, stores, or arithmetic
for the core computation
In other words: address arithmetic and loop overhead

Strategy: unroll loops

As reduction proceeds, # “active” threads decreases

When $s \leq 32$, we have only one warp left

Instructions are SIMD synchronous within a warp

That means when $s \leq 32$:

We don't need to `_syncthreads()`

We don't need “if ($tid < s$)” because it doesn't save any work

Let's unroll the last 6 iterations of the inner loop

Reduction #5: Unroll the Last Warp

```
device_void warpReduce(volatile int* sdata, int tid)
{ sdata[tid] += sdata[tid + 32];
  sdata[tid] += sdata[tid + 16];
  sdata[tid] += sdata[tid + 8];
  sdata[tid] += sdata[tid + 4];
  sdata[tid] += sdata[tid + 2];
  sdata[tid] += sdata[tid + 1];
}
```

IMPORTANT:
For this to be correct,
we must use the
“volatile” keyword!

```
// later...
for (unsigned int s=blockDim.x/2; s>32; s>>=1) {
    if (tid < s)
        sdata[tid] += sdata[tid + s];
    syncthreads();
}

if (tid < 32) warpReduce(sdata, tid);
```

Note: This saves useless work in *all* warps, not just the last one!

Without unrolling, all warps execute every iteration of the for loop and if statement

	Time (2^{22} ints)	Bandwidth	Step Speedup	Cumulative Speedup
Kernel 1: interleaved addressing with divergent branching	8.054 ms	2.083 GB/s		
Kernel 2: interleaved addressing with bank conflicts	3.456 ms	4.854 GB/s	2.33x	2.33x
Kernel 3: sequential addressing	1.722 ms	9.741 GB/s	2.01x	4.68x
Kernel 4: first add during global load	0.965 ms	17.377 GB/s	1.78x	8.34x
Kernel 5: unroll last warp	0.536 ms	31.289 GB/s	1.8x	15.01x

Complete Unrolling

If we knew the number of iterations at compile time, we could completely unroll the reduction

Luckily, the block size is limited by the GPU to 1024 threads (Tesla V100)

Also, we are sticking to power-of-2 block sizes

So we can easily unroll for a fixed block size

But we need to be generic – how can we unroll for block sizes that we don't know at compile time?

Templates to the rescue!

CUDA supports C++ template parameters on device and host functions

Unrolling with Templates

Specify block size as a function template parameter:

```
template <unsigned int blockSize>  
_global_ void reduce5(int *g_idata, int *g_odata)
```

Reduction #6: Completely Unrolled

```

Template <unsigned int blockSize>
device_ void warpReduce(volatile int* sdata, int
tid) { if (blockSize >= 64) sdata[tid] += sdata[tid +
32];
if (blockSize >= 32) sdata[tid] += sdata[tid + 16];
if (blockSize >= 16) sdata[tid] += sdata[tid + 8];
if (blockSize >= 8) sdata[tid] += sdata[tid + 4];
if (blockSize >= 4) sdata[tid] += sdata[tid + 2];
if (blockSize >= 2) sdata[tid] += sdata[tid + 1];
} if (blockSize >=

```

```

if (blockSize >= 512) {
    if (tid < 256) { sdata[tid] += sdata[tid + 256]; }_
    syncthreads(); } if (blockSize >= 256) {
    if (tid < 128) { sdata[tid] += sdata[tid + 128]; }_
    syncthreads(); } if (blockSize >= 128) {
    if (tid < 64) { sdata[tid] += sdata[tid + 64]; }_
    syncthreads(); }

```

```

if (tid < 32) warpReduce<blockSize>(sdata, tid);

```

Note: all code in **RED** will be evaluated at compile time.

Results in a very efficient inner loop!

Invoking Template Kernels

Don't we still need block size at compile time?

Nope, just a switch statement for 10 possible block sizes:

```
switch (threads)
{
    case 512:
        reduce5<512><<< dimGrid, dimBlock, smemSize >>>(d_idata, d_odata); break;
    case 256:
        reduce5<256><<< dimGrid, dimBlock, smemSize >>>(d_idata, d_odata); break;
    case 128:
        reduce5<128><<< dimGrid, dimBlock, smemSize >>>(d_idata, d_odata); break;
    case 64:
        reduce5< 64><<< dimGrid, dimBlock, smemSize >>>(d_idata, d_odata); break;
    case 32:
        reduce5< 32><<< dimGrid, dimBlock, smemSize >>>(d_idata, d_odata); break;
    case 16:
        reduce5< 16><<< dimGrid, dimBlock, smemSize >>>(d_idata, d_odata); break;
    case 8:
        reduce5< 8><<< dimGrid, dimBlock, smemSize >>>(d_idata, d_odata); break;
    case 4:
        reduce5< 4><<< dimGrid, dimBlock, smemSize >>>(d_idata, d_odata); break;
    case 2:
        reduce5< 2><<< dimGrid, dimBlock, smemSize >>>(d_idata, d_odata); break;
    case 1:
        reduce5< 1><<< dimGrid, dimBlock, smemSize >>>(d_idata, d_odata); break;
}
```

	Time (2^{22} ints)	Bandwidth	Step Speedup	Cumulative Speedup
Kernel 1: interleaved addressing with divergent branching	8.054 ms	2.083 GB/s		
Kernel 2: interleaved addressing with bank conflicts	3.456 ms	4.854 GB/s	2.33x	2.33x
Kernel 3: sequential addressing	1.722 ms	9.741 GB/s	2.01x	4.68x
Kernel 4: first add during global load	0.965 ms	17.377 GB/s	1.78x	8.34x
Kernel 5: unroll last warp	0.536 ms	31.289 GB/s	1.8x	15.01x
Kernel 6: completely unrolled	0.381 ms	43.996 GB/s	1.41x	21.16x

Replace load and add of two elements:

```
unsigned int tid = threadIdx.x;  
unsigned int i = blockIdx.x*(blockDim.x*2) + threadIdx.x;  
sdata[tid] = g_idata[i] + g_idata[i+blockDim.x];  
syncthreads();
```

With a while loop to add as many as necessary in a CUDA thread:

```
unsigned int tid = threadIdx.x;  
unsigned int i = blockIdx.x*(blockSize*2) + threadIdx.x;  
unsigned int gridSize = blockSize*2*gridDim.x;  
sdata[tid] = 0;
```

```
while (i < n) {  
    sdata[tid] += g_idata[i] + g_idata[i+blockSize];  
    i += gridSize;  
}  
syncthreads();
```

Reduction #7: More Grids

Replace load and add of two elements:

```
unsigned int tid = threadIdx.x;  
unsigned int i = blockIdx.x*(blockDim.x*2) + threadIdx.x;  
sdata[tid] = g_idata[i] + g_idata[i+blockDim.x];  
syncthreads();
```

With a while loop to add as many as necessary in a CUDA thread:

```
unsigned int tid = threadIdx.x;  
        unsigned int i = blockIdx.x*(blockSize*2) + threadIdx.x;  
unsigned int gridSize = blockSize*2*gridDim.x;  
sdata[tid] = 0;  
  
while (i < n) {  
    sdata[tid] += g_idata[i] + g_idata[i+blockSize];  
    i += gridSize;  
}  
syncthreads();
```

	Time (2^{22} ints)	Bandwidth	Step Speedup	Cumulative Speedup
Kernel 1: interleaved addressing with divergent branching	8.054 ms	2.083 GB/s		
Kernel 2: interleaved addressing with bank conflicts	3.456 ms	4.854 GB/s	2.33x	2.33x
Kernel 3: sequential addressing	1.722 ms	9.741 GB/s	2.01x	4.68x
Kernel 4: first add during global load	0.965 ms	17.377 GB/s	1.78x	8.34x
Kernel 5: unroll last warp	0.536 ms	31.289 GB/s	1.8x	15.01x
Kernel 6: completely unrolled	0.381 ms	43.996 GB/s	1.41x	21.16x
Kernel 7: multiple elements per thread	0.268 ms	62.671 GB/s	1.42x	30.04x

Kernel 7 on 32M elements: 73 GB/s!



```
template <unsigned int blockSize>
```

```
__device__ void warpReduce(volatile int *sdata, unsigned int tid) {
    if (blockSize >= 64) sdata[tid] += sdata[tid + 32];
    if (blockSize >= 32) sdata[tid] += sdata[tid + 16];
    if (blockSize >= 16) sdata[tid] += sdata[tid + 8];
    if (blockSize >= 8) sdata[tid] += sdata[tid + 4];
    if (blockSize >= 4) sdata[tid] += sdata[tid + 2];
    if (blockSize >= 2) sdata[tid] += sdata[tid + 1];
} (blockSize >=
```

Final Optimized Kernel

```
template <unsigned int blockSize>
```


```
__global__ void reduce6(int *g_idata, int *g_odata, unsigned int n)
{ __shared__ int sdata[BLOCK_SIZE];
  unsigned int tid = threadIdx.x;
  unsigned int i = blockIdx.x*(blockSize*2) + tid;
  unsigned int gridSize = blockSize*2*gridDim.x;
  sdata[tid] = 0;
```

```
while (i < n) { sdata[tid] += g_idata[i] +
g_idata[i+blockSize]; i += gridSize; }
syncthreads();
```

```
if (blockSize >= 512) { if (tid < 256) { sdata[tid] += sdata[tid + 256]; }
syncthreads(); }
if (blockSize >= 128) { if (tid < 64) { sdata[tid] += sdata[tid + 128]; }
syncthreads(); }
if (blockSize >= 64) { if (tid < 32) { sdata[tid] += sdata[tid + 64]; }
syncthreads(); }
if (tid < 32) warpReduce(sdata, tid);
if (tid == 0) g_odata[blockIdx.x] = sdata[0];
```

```
}
```

Performance Comparison



Interesting observation:

Algorithmic optimizations

Changes to addressing, algorithm cascading
11.84x speedup, combined!

Code optimizations

Loop unrolling
2.54x speedup, combined

Understand CUDA performance characteristics

- Memory coalescing

- Divergent branching

- Bank conflicts

- Latency hiding

Use peak performance metrics to guide optimization

Understand parallel algorithm complexity theory

Know how to identify type of bottleneck

- e.g. memory, core computation, or instruction overhead

Optimize your algorithm, *then* unroll loops

Use template parameters to generate optimal code

- Using Managed Memory
- Concurrency (streams, copy/compute overlap, multi-GPU)
- Analysis Driven Optimization
- Cooperative Groups

- Parallel reduction:
 - <https://developer.download.nvidia.com/assets/cuda/files/reduction.pdf>
- Warp-shuffle and reduction:
 - <https://devblogs.nvidia.com/faster-parallel-reductions-kepler/>
- CUDA Cooperative Groups:
 - <https://devblogs.nvidia.com/cooperative-groups/>
- Grid-stride loops:
 - <https://devblogs.nvidia.com/cuda-pro-tip-write-flexible-kernels-grid-stride-loops/>
- Floating point:
 - <https://developer.nvidia.com/sites/default/files/akamai/cuda/files/NVIDIA-CUDA-Floating-Point.pdf>
- CUDA Sample Codes:
 - Reduction, threadFenceReduction, reductionMultiBlockCG